

Chapter 7

MLR Model Evaluation

Chanwoo Yoo, Division of Advanced Engineering,
Korea National Open University

This work is a derivative of 'Regression Methods' by Iain Pardoe, Laura Simon and Derek Young, used under CC BY-NC.

Contents

1. Three Types of Hypotheses
2. The General Linear F-Test

1. Three Types of Hypotheses

1. Three Types of Hypotheses

- A hypothesis test for testing that one slope parameter is 0
- A hypothesis test for testing that all of the slope parameters are 0
- A hypothesis test for testing that a subset — more than one, but not all — of the slope parameters are 0

2. Example: Heart attacks in Rabbits

- Does the mean size of the infarcted area differ among the three treatment groups — no cooling, early cooling, and late cooling — when controlling for the size of the region at risk for infarction?

2. Example: Heart attacks in Rabbits

- Rabbits whose hearts were cooled to 6° C within 5 minutes of the blocked artery ("early cooling")
- Rabbits whose hearts were cooled to 6° C within 25 minutes of the blocked artery ("late cooling")
- Rabbits whose hearts were not cooled at all ("no cooling")

3. Data: Cool Hearts

- Data: [Cool Hearts](#)
 - y_i (Inf.): the size of the infarcted area (in grams) of rabbit i
 - x_{i1} (Area): the size of the region at risk (in grams) of rabbit i
 - x_{i2} (X2): 1 if early cooling of rabbit i , 0 if not
 - x_{i3} (X3): 1 if late cooling of rabbit i , 0 if not
 - $y_i = (\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \beta_3 x_{i3}) + \epsilon_i$

4. Code: Cool Hearts

```
coolhearts <- read.table("coolhearts.txt", header=T)  
attach(coolhearts)
```

```
model.1 <- lm(Inf. ~ Area + X2 + X3)  
summary(model.1)  
anova(model.1)
```


5. Results: Cool Hearts

```
> summary(model.1)
```

```
...
```

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-0.13454	0.10402	-1.293	0.206459	
Area	0.61265	0.10705	5.723	3.87e-06	***
X2	-0.24348	0.06229	-3.909	0.000536	***
X3	-0.06566	0.06507	-1.009	0.321602	

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

5. Results: Cool Hearts

```
> summary(model.1)
```

```
...
```

```
Residual standard error: 0.1395 on 28 degrees of freedom
```

```
Multiple R-squared: 0.6377, Adjusted R-squared: 0.5989
```

```
F-statistic: 16.43 on 3 and 28 DF, p-value: 2.363e-06
```

5. Results: Cool Hearts

```
> anova(model.1)
```

Analysis of Variance Table

Response: Inf.


	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
Area	1	0.62492	0.62492	32.1115	4.504e-06	***
X2	1	0.31453	0.31453	16.1622	0.000398	***
X3	1	0.01981	0.01981	1.0181	0.321602	
Residuals	28	0.54491	0.01946			

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

6. Code: Cool Hearts


```
plot(Area, Inf., type="n", ylim=c(-0.2, 1),  
     xlab="Size of area at risk (grams)",  
     ylab="Size of infarcted area (grams)")  
for (i in 1:32) points(Area[i], Inf.[i], pch=Group[i],  
col=Group[i])  
for (i in 1:3) lines(Area[Group==i], fitted(model.1)[Group==i],  
lty=i, col=i)  
legend("topleft", legend=c("no cooling",  
                           "late cooling",  
                           "early cooling"),  
      col=3:1, pch=3:1, inset=0.02)
```

7. Results: Cool Hearts




- As the size of the area at risk increases, the size of the infarcted area also tends to increase.

7. Results: Cool Hearts



- The plot also suggests that for this sample of 32 rabbits with a given size of area at risk, the average size of the infarcted area differs for the three experimental groups.

7. Results: Cool Hearts



- We want to be able to answer the research question for the whole population of rabbits.

8. Statistical Procedure

- To translate a research question into a statistical procedure,
 - formulate a multiple regression model
 - determine how the model helps answer the research question
 - check the model
 - perform a hypothesis test (or calculating a confidence interval)

9. Research Question I

- Is a regression model containing at least one predictor useful in predicting the size of the infarct?
 - $H_0: \beta_1 = \beta_2 = \beta_3 = 0$
 - $H_A: \text{At least one } \beta_i \neq 0 \text{ (for } i = 1, 2, 3)$

10. Research Question II

- Is the size of the infarct significantly (linearly) related to the area of the region at risk?
 - $H_0: \beta_1 = 0$
 - $H_A: \beta_1 \neq 0$

11. Research Question III

- Is the size of the infarct area significantly (linearly) related to the type of treatment after controlling for the size of the region at risk for infarction?
 - $H_0: \beta_2 = \beta_3 = 0$
 - H_A : At least one $\beta_i \neq 0$ (for $i = 2,3$)

2. The General Linear F-Test

1. The General Linear F-Test


- The "general linear F-test" involves three basic steps
 - Define a larger full model.
 - Define a smaller reduced model.
 - Use an F-statistic to decide whether or not to reject the smaller reduced model in favor of the larger full model.

2. The Full Model

- The "full model", which is also sometimes referred to as the "unrestricted model," is the model thought to be most appropriate for the data. For simple linear regression, the full model is

$$y_i = (\beta_0 + \beta_1 x_{i1}) + \epsilon_i$$

2. The Full Model




3. The Reduced Model

- The "reduced model," which is sometimes also referred to as the "restricted model," is the model described by the null hypothesis H_0 . For simple linear regression, a common null hypothesis is $H_0: \beta_1 = 0$. In this case, the reduced model is obtained by "zeroing-out" the slope β_1 that appears in the full model. That is, the reduced model is

$$y_i = \beta_0 + \epsilon_i$$

3. The Reduced Model



4. F-Statistic Test

- How do we decide if the reduced model or the full model does a better job of describing the trend in the data when it can't be determined by simply looking at a plot? What we need to do is to quantify how much error remains after fitting each of the two models to our data. That is, we take the general linear F-test approach.

4. F-Statistic Test

- "Fit the full model" to the data.
 - Obtain the least squares estimates of β_0 and β_1 .
 - Determine the error sum of squares, which we denote "SSE(F)."
- "Fit the reduced model" to the data.
 - Obtain the least squares estimate of β_0 .
 - Determine the error sum of squares, which we denote "SSE(R)."

4. F-Statistic Test

- The general linear F-statistic
 - $F^* = \left(\frac{SSE(R) - SSE(F)}{df_R - df_F} \right) \div \left(\frac{SSE(F)}{df_F} \right)$
 - df_R, df_F : The degrees of freedom
- We use the general linear F-statistic to decide whether or not
 - to reject the null hypothesis H_0 : The reduced model
 - in favor of the alternative hypothesis H_A : The full model

5. The test applied to the simple linear regression model

- $H_0: y_i = \beta_0 + \epsilon_i$
- $H_A: y_i = (\beta_0 + \beta_1 x_{i1}) + \epsilon_i$

5. The test applied to the simple linear regression model

$$\blacksquare F^* = \left(\frac{SSE(R) - SSE(F)}{df_R - df_F} \right) \div \left(\frac{SSE(F)}{df_F} \right)$$

$$\bullet df_R = n - 1, df_F = n - 2$$

$$\bullet SSE(R) = SSTO, SSE(F) = SSE$$

$$\blacksquare F^* = \left(\frac{SSE(R) - SSE(F)}{df_R - df_F} \right) \div \left(\frac{SSE(F)}{df_F} \right) = \left(\frac{SSTO - SSE}{(n-1) - (n-2)} \right) \div \left(\frac{SSE}{n-2} \right) = \frac{MSR}{MSE}$$

6. Example

```
skincancer <- read.table("skincancer.txt", header=T)
attach(skincancer)
model <- lm(Mort ~ Lat)

summary(model)
anova(model)
```

6. Example

```
> anova(model)
```

Analysis of Variance Table

Response: Mort

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Lat	1	36464	36464	99.797	3.309e-13 ***
Residuals	47	17173	365		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

$$SSTO = 36464 + 17173 = 53637$$

$$SSE = 17173$$

$$MSR = 36464$$

$$MSE = 365$$

SSE

MSE

6. Example

- $MSR = 36464$
- $MSE = 365$
- $F^* = \frac{MSR}{MSE} = \frac{36464}{365} = 99.9$
 - 1 numerator degree of freedom and 47 denominator degree of freedom

6. Example

```
> pf(99.9, 1, 47, lower.tail = FALSE)
[1] 3.254756e-13
```

6. Example

```
> summary(model)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	389.1894	23.8123	16.34	< 2e-16	***
Lat	-5.9776	0.5984	-9.99	3.31e-13	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 19.12 on 47 degrees of freedom

Multiple R-squared: 0.6798, Adjusted R-squared: 0.673

F-statistic: 99.8 on 1 and 47 DF, p-value: 3.309e-13

6. Example

```
> anova(model)
```

Analysis of Variance Table

Response: Mort

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Lat	1	36464	36464	99.797	3.309e-13 ***
Residuals	47	17173	365		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Next

Chapter 8

MLR Model Evaluation II